



What is Statistics?


1000's of units

context: statistical consultant, hired to give advice on improving the

output variables (response, dependent variable, etc.)



input variables (predictors, features, independent variables)



Shown are **Sales** vs **TV**, **Radio**, and **Newspaper**, with a blue linear-regression line fit separately to each.

Can we predict **Sales** using these three?

Perhaps we can do better using a model

$$\text{Sales} \approx f(\text{TV}, \text{Radio}, \text{Newspaper})$$

random variable

fixed variables (not random)
H) *use or target* that we wish to predict. We generically refer to the response as Y .

TV is a *feature, or input, or predictor*; we name it X_1 .

Likewise name **Radio** as X_2 , and so on.

We can refer to the *input vector* collectively as

fixed (not random)

"data generating model":
the model that we presume
our data comes from

$$X = \begin{pmatrix} X_1 \\ X_2 \\ X_3 \end{pmatrix}$$

$Y(x)$: the random response
associated with fixed inputs
 x

Now we write our model as

$f(x)$: a systematic function
of inputs x (not random)

$$Y = f(X) + \epsilon$$

random variable modeling
error (capture things we
cannot control).
- assume $E(\epsilon)$ (i.e.,
its expectation) = 0. zero-
mean error term

where ϵ captures measurement errors and

Two types of predictions:

- point prediction
- probabilistic predictions (predicting a distribution)

What is $f(X)$

Two properties we may want to learn from a model:

- (i) prediction
- (ii) inference


- With a good f we can make predictions of Y at new points $X = x$.
- We can understand which components $X = (X_1, X_2, \dots, X_p)$ are important in which are irrelevant. e.g. **Seniority** & **Education** have a big impact on **Income** **Status** typically does not.
- Depending on the complexity of f , we can understand how each component X_j contributes to the prediction.

inference: understanding the domain problem better from data

- which predictors are associated / influential for the response?

- what is the relationship between a predictor and a response?

- is this relationship linear? or is it more complex (nonlinear)?



Is there an ideal $f(X)$? In particular, what is a good value for $f(X)$ at any selected value of X , say $X = 4$? There can be many Y values at $X = 4$. A good value is

Let's start by assuming we know the generative model (with certainty):

$$E(Y | X = x) = f(x) + \epsilon = 4.$$

This

- we know what the function f is precisely
- we also know what the random variable ϵ is (e.g., its pdf or pmf)

function.

The regression function $f(x)$

- Is also defined for vector X ; e.g.

$$f(x) = f(x_1, x_2, x_3) = E(Y|X_1 = x_1, X_2 = x_2, X_3 = x_3)$$

The regression

mean-squared prediction error (MSPE):
- one choice of loss function

- Is also defined for vector X ; e.g.
$$f(x) = f(x_1, x_2, x_3) = E(Y|X_1 = x_1, X_2 = x_2, X_3 = x_3)$$
- Is the *ideal* or *optimal* predictor of Y with regard to mean-squared prediction error: $f(x) = E(Y|X = x)$ is the function that minimizes $E[(Y - g(X))^2|X = x]$ over all functions g at all points $X = x$.

The regression function $f(x)$

- Is also defined for vector X ; e.g.
$$f(x) = f(x_1, x_2, x_3) = E(Y|X_1 = x_1, X_2 = x_2, X_3 = x_3)$$
- Is the *ideal* or *optimal* predictor of Y with regard to mean-squared prediction error: $f(x) = E(Y|X = x)$ is the function that minimizes $E[(Y - g(X))^2|X = x]$ over all functions g at all points $X = x$.
- $\epsilon = Y - f(x)$ is the *irreducible* error — i.e. even if we knew $f(x)$, we would still make errors in prediction, since at each $X = x$ there is typically a distribution of possible Y values.

what happens to reducible error as we collect more and more data?

- model misspecification: the assumed model for estimating $f(x)$ is not the true regression model (e.g., last example, if we used a linear regression model to estimate $f(x)$)
- when there is model misspecification, then reducible error will decrease as we collect more data, but NOT go to zero (systematic bias)
- assuming no model misspecification, reducible error $\rightarrow 0$ as we collect more data

regression function $f(x)$

predictor X ; e.g.

$$f(x) = E(Y|X_1 = x_1, X_2 = x_2, X_3 = x_3)$$

best predictor of Y with regard to

prediction error: $f(x) = E(Y|X = x)$ is the function that minimizes $E[(Y - g(X))^2|X = x]$ over all functions g that depend only on $X = x$.

irreducible error —

cannot be reduced by our choice of predictor g

$f(x)$, we would still make errors in prediction, since at each $X = x$ there is typically a distribution of possible Y values.

- For any estimate $\hat{f}(x)$ of $f(x)$, we have

$$E[(Y - \hat{f}(X))^2|X = x] = \underbrace{[f(x) - \hat{f}(x)]^2}_{\text{Reducible}} + \underbrace{\text{Var}(\epsilon)}_{\text{Irreducible}}$$

See board for derivation

One strategy: K-nearest neighbor prediction at $x=4$

- take the closest K (e.g., 10) data points to $x=4$
- set our predictor of $f(4)$ (denote as $\hat{f}(4)$) as simply the average of the K outputs from the data

How to estimate f

Have few if any data points with $X = 4$

not compute $E(Y|X = x)$!


- Relax the definition and let

$$\hat{f}(x) = \text{Ave}(Y|X \in \mathcal{N}(x))$$

where $\mathcal{N}(x)$ is some *neighborhood*

K-nearest neighbors (KNNs) are a popular class of non-parametric model for estimating $f(x)$

green line = prediction of $f(x)$





- Nearest neighbor averaging can be pretty good for small p — i.e. $p \leq 4$ and large-ish N .
- We will discuss smoother versions, such as kernel and spline smoothing later in the course.

- Nearest neighbor averaging can be pretty good for small p — i.e. $p \leq 4$ and large-ish N .
- We will discuss smoother versions, such as kernel and spline smoothing later in the course.
- Nearest neighbor methods can be *lousy* when p is large. Reason: the *curse of dimensionality*. Nearest neighbors tend to be far away in high dimensions.
 - We need to get a reasonable fraction of the N values of y_i to average to bring the variance down—e.g. 10%.
 - A 10% neighborhood in high dimensions need no longer be local, so we lose the spirit of estimating $E(Y|X = x)$ by local averaging.

The curse of dimensionality

10% Neighborhood



Parametric model:

- assume a model on f that depends on a finite number of parameters (β)
- estimate the parameters β to obtain a good predictor \hat{f} for the true regression function $f(x)$

Linear and structured models


Linear model is an important example of a parametric

$$f_L(X) = \beta_0 + \beta_1 X_1 + \beta_2 X_2$$


i.e., the true regression function $f(x)$ is almost never a linear function of x

- A linear model is specified in terms of $p + 1$ parameters $\beta_0, \beta_1, \dots, \beta_p$.
- We estimate the parameters by fitting the model to training data.
- Although it is *almost never correct*, a linear model often serves as a good and interpretable approximation to the unknown true function $f(X)$.

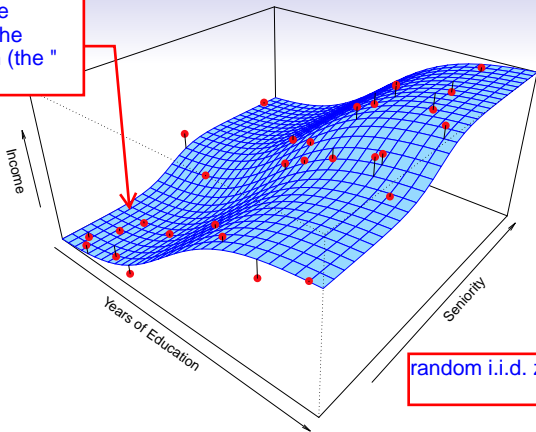
A linear model $\hat{f}_L(X) = \hat{\beta}_0 + \hat{\beta}_1 X$ gives a reasonable fit here



A quadratic model $\hat{f}_Q(X) = \hat{\beta}_0 + \hat{\beta}_1 X + \hat{\beta}_2 X^2$ fits slightly better.



we're simulating the training data from the regression function (the "truth")

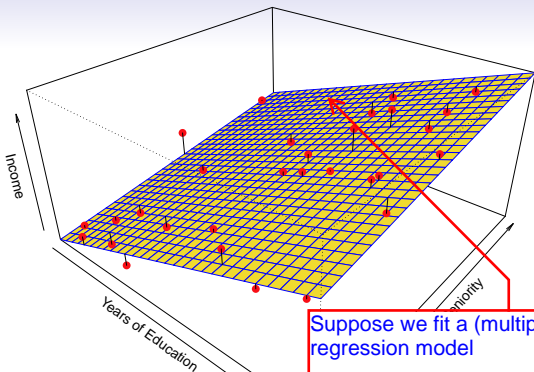


random i.i.d. zero-mean

Simulated example. Red points are simulated values for **income** from the model

$$\text{income} = f(\text{education}, \text{seniority}) + \epsilon$$

f is the blue surface.



Suppose we fit a (multiple) linear regression model


Observations:

- overall: not too bad of a fit to the data
- but there is some evidence of systematic "bias" in the fitted model:
 - some of the fitted errors are mostly positive in certain regions, and others mostly negative
 - assuming that there is model misspecification, reducible error will not go to 0 as sample size increases

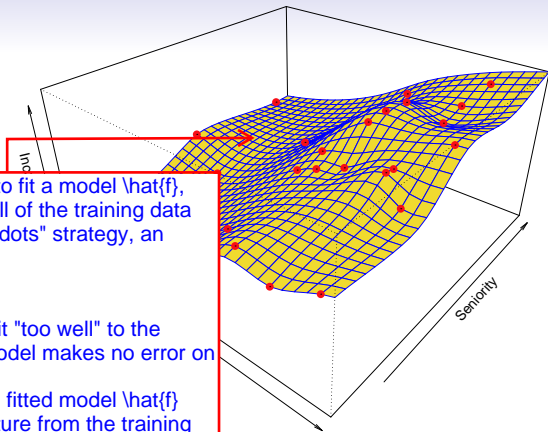
Linear regression model fit to the si

$$\hat{f}_L(\text{education}, \text{seniority}) = \hat{\beta}_0 + \hat{\beta}_1$$

Non-parametric modeling:
- much less assumptions
on the form of f
- e.g., KNN, thin-plate
splines



More flexible regression model $\hat{f}_S(\text{education}, \text{seniority})$ fit to the simulated data. Here we use a technique called a *thin-plate spline* to fit a flexible surface. We control the roughness of the fit (chapter 7).



Suppose we choose to fit a model \hat{f} , which goes through all of the training data points - "connect-the-dots" strategy, an interpolator

Observations:

- predictor seems to fit "too well" to the training data, fitted model makes no error on training data
- overfitting: when the fitted model \hat{f} learns incorrect structure from the training data noise
- results in an overly complex / overly flexible model
- an overfit model has high VARIANCE:
 - if we draw a new training dataset from the generative model, and we fit a new predictor \hat{f}^* , then \hat{f}^* will likely be very different from \hat{f}
 - think connect-the-dots

gression model

fit to the simulated data. Here the
s on the training data! Also known

Some trade-offs


- Prediction accuracy versus interpretability.
 - Linear models are easy to interpret; thin-plate splines are not.

Some trade-offs

- Prediction accuracy versus interpretability.
 - Linear models are easy to interpret; thin-plate splines are not.
- Good fit versus over-fit or under-fit.
 - How do we know when the fit is just right?

Some trade-offs

- Prediction accuracy versus interpretability.
 - Linear models are easy to interpret; thin-plate splines are not.
- Good fit versus over-fit or under-fit.
 - How do we know when the fit is just right?
- Parsimony versus black-box.
 - We often prefer a simpler model involving fewer variables over a black-box predictor involving them all.



Assessing Model Accuracy

Suppose we fit a model $\hat{f}(x)$ to some training data $\text{Tr} = \{x_i, y_i\}_1^N$, and we wish to see how well it performs.

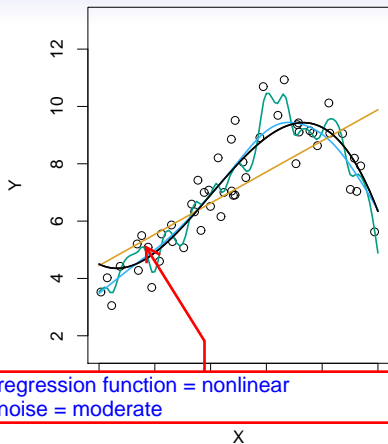
- We could compute the average squared prediction error over Tr :

$$\text{MSE}_{\text{Tr}} = \text{Ave}_{i \in \text{Tr}} [y_i - \hat{f}(x_i)]^2$$

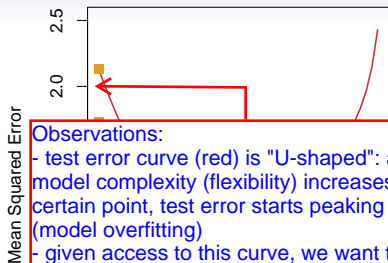
This may be biased toward more overfit models.

- Instead we should, if possible, compute it using fresh *test* data $\text{Te} = \{x_i, y_i\}_1^M$:

$$\text{MSE}_{\text{Te}} = \text{Ave}_{i \in \text{Te}} [y_i - \hat{f}(x_i)]^2$$



regression function = nonlinear
noise = moderate



Observations:


- test error curve (red) is "U-shaped": as model complexity (flexibility) increases past a certain point, test error starts peaking up (model overfitting)
- given access to this curve, we want to be at the "optimal" model complexity to minimize MSPE (blue model)
- training MSPE curve (gray): always decreases as model complexity increases, hence it's not a good criterion for prediction
- testing MSPE generally higher than training MSPE. why? model is fitted to minimize training error, and not testing

Black curve is truth. Red curve on right is MSE_{Te} , grey curve is MSE_{Tr} . Orange, blue and green curves/squares correspond to fits of different flexibility.

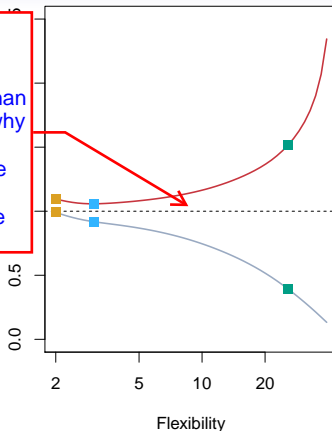
Observations:

- testing MSPE is higher than training MSPE
- "U-shaped" test error
- the optimal model flexibility is much lower than the optimal flexibility in the earlier example. why?


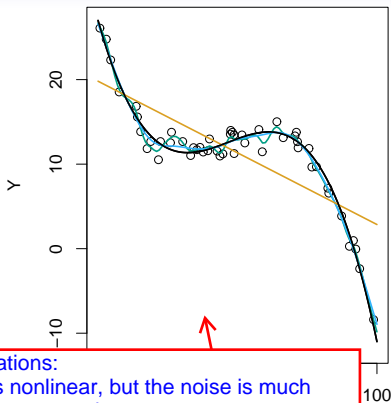
- variance curve remains largely the same (noise level is similar)
- bias curve is very different since the true regression function is nearly linear



regression function = almost linear
noise = moderate



Here the truth is smoother, so the smoother fit and linear model do really well.



Observations:

- truth is nonlinear, but the noise is much lower than example 1
- there is much less danger in overfitting the model
 - bias curve remains largely the same as example 1 (similar degrees of nonlinearity)
 - variance curve increases much more slowly, since with less noise there is less risk of overfitting
- the optimal model for prediction is much more complex

Since the noise is low, so the more flexible fits

Bias-Va

Suppose we have fit a model
let (x_0, y_0) be a test observ
the true model is $Y = f(x)$
then

$$E \left(y_0 - \hat{f}(x_0) \right)^2 = \text{Var}$$

The expectation averages over
the variability in Tr. Note

Typically as the *flexibility* of \hat{f} increases, its variance increases,
and its bias decreases. So choosing the flexibility based on
average test error amounts to a *bias-variance trade-off*.

derivation on board

- LHS: testing MSPE which we wish to minimize
- RHS: reducible + irreducible error (from earlier)

The reducible error has two components: bias & variance

BIAS:




- "error from approximating a complex regression function with a simpler model", i.e., model misspecification
- linear models are often biased, unless the true regression function is linear

VARIANCE:

- "how much the predictor \hat{f} changes when we estimate this with a new training dataset"
- linear models have little variance: smaller # of parameters to fit with the data
- flexible models have high variance (connect-the-dots)

TRADE-OFF: find an optimal compromise between bias and variance for fitting \hat{f}

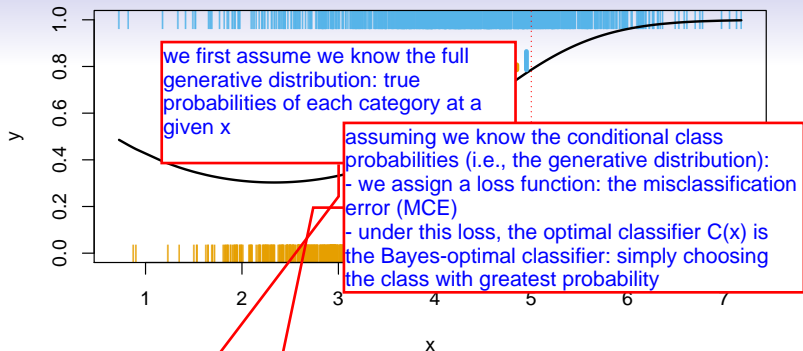
Bias-variance trade-off for the three examples



Classification Problems

Here the response variable Y is *qualitative* — e.g. email is one of $\mathcal{C} = (\text{spam}, \text{ham})$ (ham =good email), digit class is one of $\mathcal{C} = \{0, 1, \dots, 9\}$. Our goals are to:

- Build a classifier $C(X)$ that assigns a class label from \mathcal{C} to a future unlabeled observation X .
- Assess the uncertainty in each classification
- Understand the roles of the different predictors among $X = (X_1, X_2, \dots, X_p)$.



Is there an ideal $C(X)$? Suppose the K elements in \mathcal{C} are numbered $1, 2, \dots, K$. Let

$$p_k(x) = \Pr(Y = k | X = x), \quad k = 1, 2, \dots, K.$$

These are the *conditional class probabilities* at x ; e.g. see little barplot at $x = 5$. Then the *Bayes optimal* classifier at x is


$$C(x) = j \text{ if } p_j(x) = \max\{p_1(x), p_2(x), \dots, p_K(x)\}$$

Similar to the regression problem: we don't know what the conditional class probabilities are in practice!

- we want to estimate these probabilities via a model trained from data
- then use the trained model to build the Bayes optimal classifier

- an alternate (parametric) model for the conditional class probabilities would be logistic regression.

- the same principles underlying the bias-variance trade-off (in choosing optimal model complexity) applies here



K-nearest-neighbor (KNN) classifier:

- e.g., let's build the trained classifier at $x=5$ with $K=10$ neighbors
- we see that there are 4 0's and 6 1's of the $K=10$ nearest neighbors
- using this, our estimate of $p_1(x) = 60\%$, of $p_0(x) = 40\%$
- using these estimated probabilities, since our estimate of $p_1(x) > p_0(x)$, our estimated Bayes optimal classifier is class 1

Nearest-neighbor and
Also breaks down a
 $\hat{C}(x)$ is less than or

ect on

Classification: some details

- Typically we measure the performance of $\hat{C}(x)$ using the misclassification error rate:

$$\text{Err}_{\text{Te}} = \text{Ave}_{i \in \text{Te}} I[y_i \neq \hat{C}(x_i)]$$

- The Bayes classifier (using the true $p_k(x)$) has smallest error (in the population).

Classification: some details


- Typically we measure the performance of $\hat{C}(x)$ using the misclassification error rate:

$$\text{Err}_{\text{T}_e} = \text{Ave}_{i \in \text{T}_e} I[y_i \neq \hat{C}(x_i)]$$

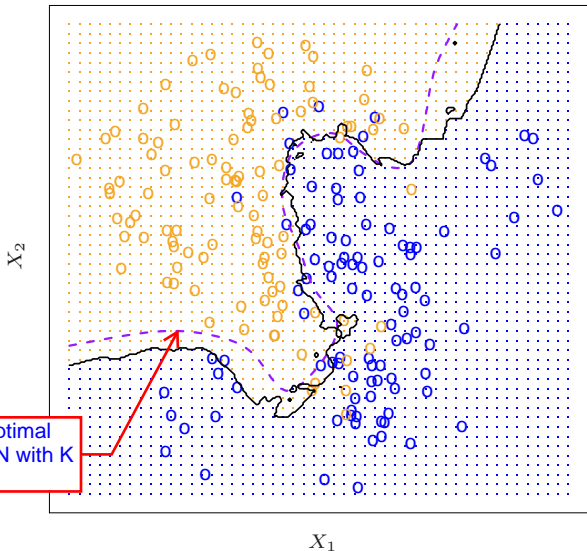
- The Bayes classifier (using the true $p_k(x)$) has smallest error (in the population).
- Support-vector machines build structured models for $C(x)$.
- We will also build structured models for representing the $p_k(x)$. e.g. Logistic regression, generalized additive models.

Example: K-nearest neighbors in two dimensions

purple line: true Bayes
optimal classifier
- the optimal classifier if we
know the true conditional
class probabilities




KNN: K=10



estimated Bayes optimal
classifier using KNN with K
= 10 neighbors


black line: trained classifier
using KNN with $K = 1$
- not too close to the true
optimal classifier (purple)
- overfitting to the training
data (variance)

$K=1$



black line: trained classifier
using KNN with $K=100$
- not close at all to the true
optimal classifier (purple)
- underfitting to the training
data (bias)

$K=100$



The number of neighbors K controls the complexity of
our classifier: bias-variance trade-off

