


Support Vector Machines



a popular classification
method in the computer
science community:
- one of the best black box
classifier

Here we approach the two-class problem in a direct way:

We try and find a plane that separates the classes in feature space.

If we cannot, we get creative in two ways:

- We soften what we mean by “separates”, and
- We enrich and enlarge the feature space so that separation is possible.

What is a Hyperplane?

- A hyperplane in p dimensions is a flat affine subspace of dimension $p - 1$.
- In general the equation for a hyperplane has the form

$$\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p = 0$$

- In $p = 2$ dimensions a hyperplane is a line.
- If $\beta_0 = 0$, the hyperplane goes through the origin, otherwise not.
- The vector $\beta = (\beta_1, \beta_2, \dots, \beta_p)$ is called the normal vector — it points in a direction orthogonal to the surface of a hyperplane.

Hyperplane in 2 Dimensions


the normal vector for the hyperplane

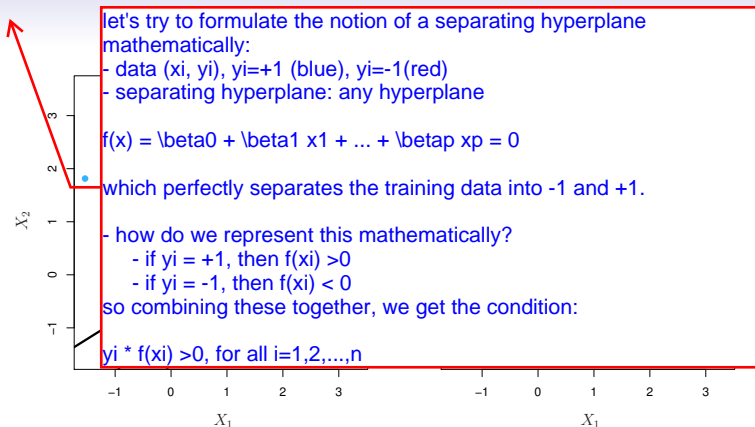
$$\beta_0 + \beta_1 x_1 + \beta_2 x_2 = 0$$

is the vector (β_1, β_2)

we can then split the domain into two regions:

- $\beta_1 x_1 + \beta_2 x_2 - 6 < 0$ (bottom of the line)
- $\beta_1 x_1 + \beta_2 x_2 - 6 > 0$ (top of the line)





- If $f(X) = \beta_0 + \beta_1 X_1 + \dots + \beta_p X_p$, then $f(X) > 0$ for points on one side of the hyperplane, and $f(X) < 0$ for points on the other.
- If we code the colored points as $Y_i = +1$ for blue, say, and $Y_i = -1$ for mauve, then if $Y_i \cdot f(X_i) > 0$ for all i , $f(X) = 0$ defines a *separating hyperplane*.

Maximal Margin Classifier

Intuition behind Maximal Margin Classifier (MMC):

- suppose I give you a separating hyperplane $f(x) = 0$
- claim: for any point x , $|f(x)|$ measures the closest distance from x to the hyperplane (under the constraint of $\sum \beta_j^2 = 1$)

- can show using linear algebra

- the condition $y_i f(x_i) \geq M$:

- provides a margin for the separation condition
- ensures each point is at least distance M from the separating hyperplane $f(x) = 0$
- margin M provides a measure of "confidence" that an observation is correctly classified (larger $M \Rightarrow$ greater confidence)

- we wish to find a hyperplane (by finding β_0, \dots, β_p) which maximizes this margin M .

- support vectors:

and the one that makes the
to classes.

trained optimization problem

maximize M
 $\beta_0, \beta_1, \dots, \beta_p$


subject to $\sum_{j=1}^p \beta_j^2 = 1,$

$y_i(\beta_0 + \beta_1 x_{i1} + \dots + \beta_p x_{ip}) \geq M$

for all $i = 1, \dots, N$.

Maximal Margin Classifier

Among all separating hyperplanes, find the one that makes the biggest gap or margin between the two classes.



Constrained optimization problem

$$\text{maximize } M$$

$$\beta_0, \beta_1, \dots, \beta_p$$

$$\text{subject to } \sum_{j=1}^p \beta_j^2 = 1,$$

$$y_i(\beta_0 + \beta_1 x_{i1} + \dots + \beta_p x_{ip}) \geq M$$

for all $i = 1, \dots, N$.

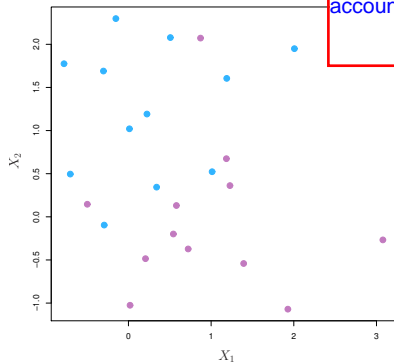


This can be rephrased as a convex quadratic program, and solved efficiently. The function `svm()` in package `e1071` solves this problem efficiently

Non-separable Data

What if the data does not have a separating hyperplane?

- MMC optimization has no solution for $M > 0$
- this is often the case for lower-dimensional data ($N > p$)
- we need to extend the notion of "separation" to account for this setting





not separable by a linear boundary.

This is often the case, unless $N < p$.

Noisy Data

sensitivity: if a new observation is added to the training point, MMC can change dramatically

- high variance, overfitting



Some
poor

noisy data: if the true class probability $P(x=-1) = 0.49$, then the optimal classifier is $+1$ (red), but we can easily get a label of -1 (blue) in the training data


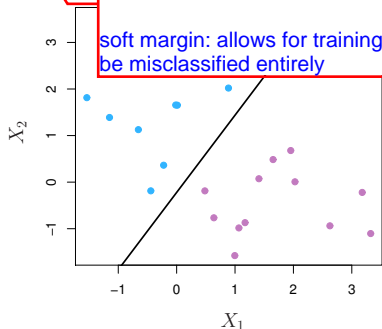
- by perfectly classifying the data point, MMC would give an incorrect classification
- overfitting, high variance of MMC

noisy. This can lead to a classifier.

Noisy Data

hard margin: all training data are classified correctly and satisfies margin (distance of M away from the hyperplane)
=> MMC

soft margin: allows for training data to violate the margin or be misclassified entirely





Sometimes the data are separable, but noisy. This can lead to a poor solution for the maximal-margin classifier.

The *support vector classifier* maximizes a *soft* margin.

Support Vector Classifier

observations 1 & 8: violate margin but correctly classified

observations 11 & 12: violate margin and misclassified



Comparing to MMC, what's new here?

- ϵ_i : slack (relaxation) variables which relax the hard margin constraint
- $\epsilon_i = 0$: constraint becomes a hard margin (i-th point must be classified)


$$\underset{\beta_0, \beta_1, \dots, \beta_p, \epsilon_1, \dots, \epsilon_n}{\text{maximize}} \quad M \quad \text{subject to} \quad \sum_{j=1}^p \beta_j^2 = 1,$$

$$y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) \geq M(1 - \epsilon_i)$$

$$\epsilon_i \geq 0, \quad \sum_{i=1}^n \epsilon_i \leq C,$$

C is a regularization parameter

controls the bias-variance trade-off of the classifier




C large:

- high tolerance budget for margin violations
- more observations are allowed to violate the margin
- margin M is large

... many observations are involved in determining the optimal hyperplane (many support vectors)

classifier will have low variance but high bias




C small:


- lower tolerance budget for margin violations
- fewer observations are allowed to violate the margin
- margin M is small

... only a few observations are involved in determining the optimal hyperplane (few support vectors)

classifier will have high variance but low bias



Linear boundary can fail



Sometime a linear boundary simply won't work, no matter what value of C .

The example on the left is such a case.

What to do?

Feature Expansion

- Enlarge the space of features by including transformations; e.g. X_1^2 , X_1^3 , X_1X_2 , $X_1X_2^2$, ... Hence go from a p -dimensional space to a $M > p$ dimensional space.
- Fit a support-vector classifier in the enlarged space.
- This results in non-linear decision boundaries in the original space.

Feature Expansion

- Enlarge the space of features by including transformations; e.g. X_1^2 , X_1^3 , X_1X_2 , $X_1X_2^2$, ... Hence go from a p -dimensional space to a $M > p$ dimensional space.
- Fit a support-vector classifier in the enlarged space.
- This results in non-linear decision boundaries in the original space.

Example: Suppose we use $(X_1, X_2, X_1^2, X_2^2, X_1X_2)$ instead of just (X_1, X_2) . Then the decision boundary would be of the form

$$\beta_0 + \beta_1X_1 + \beta_2X_2 + \beta_3X_1^2 + \beta_4X_2^2 + \beta_5X_1X_2 = 0$$


This leads to nonlinear decision boundaries in the original space (quadratic conic sections).

Cubic Polynomials

Here we use a basis expansion of cubic polynomials

From 2 variables to 9

$f(x) < 0$
The classifier in the enlarged space solves the problem in the lower-dimensional space



$f(x) = 0$
- nonlinear in the original $d=2$ dimensional space
- linear in the expanded 9-dimensional space


$f(x) > 0$

Cubic Polynomials

Here we use a basis expansion of cubic polynomials

From 2 variables to 9

The support-vector classifier in the enlarged space solves the problem in the lower-dimensional space



$$\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_1^2 + \beta_4 X_2^2 + \beta_5 X_1 X_2 + \beta_6 X_1^3 + \beta_7 X_2^3 + \beta_8 X_1 X_2^2 + \beta_9 X_1^2 X_2 = 0$$

Nonlinearities and Kernels

- Polynomials (especially rather fast).

one computational bottleneck with previous strategy:

- p variables, d -th order polynomials for the classifier
- $O(d^p)$ total features to optimize in the SVC
- thus, if you want to fix highly complex nonlinear classifiers, SVC optimization can be very expensive

kernels provide a nice computational trick to bypass this issue

- There is a more elegant and controlled way to introduce nonlinearities in support-vector classifiers — through the use of *kernels*.
- Before we discuss these, we must understand the role of *inner products* in support-vector classifiers.

Inner products and support vectors

$$\langle x_i, x_{i'} \rangle = \sum_{j=1}^p x_{ij} x_{i'j} \quad \text{— inner product between vectors}$$

Inner products and support vectors

$$\langle x_i, x_{i'} \rangle = \sum_{j=1}^p x_{ij} x_{i'j} \quad \text{— inner product between vectors}$$

- The linear support vector classifier can be represented as

$$f(x) = \beta_0 + \sum_{i=1}^n \alpha_i \langle x, x_i \rangle \quad \text{— } n \text{ parameters}$$

Inner products and support vectors

$$\langle x_i, x_{i'} \rangle = \sum_{j=1}^p x_{ij} x_{i'j} \quad \text{— inner product between vectors}$$

- The linear support vector classifier can be represented as

$$f(x) = \beta_0 + \sum_{i=1}^n \alpha_i \langle x, x_i \rangle \quad \text{— } n \text{ parameters}$$

- To estimate the parameters $\alpha_1, \dots, \alpha_n$ and β_0 , all we need are the $\binom{n}{2}$ inner products $\langle x_i, x_{i'} \rangle$ between all pairs of training observations.

Inner products and support vectors

note:

- if $\alpha_i = 0$: point i does not determine or affect the SVC (point i is not a support vector)
- if α_i is nonzero: point i is a support vector
- we know there are only a

- few support vectors \Rightarrow only a few of the α 's would be non-zero
- vector classifier can be represented as

CLAIM: the optimal SVC solution $f(x)$ can be written in this form, for some choice of β_0 and $\alpha_1, \dots, \alpha_n$

$$f(x) = \beta_0 + \sum_{i=1}^n \alpha_i \langle x, x_i \rangle \quad \text{--- } n \text{ parameters}$$

- To estimate the parameters $\alpha_1, \dots, \alpha_n$ and β_0 , all we need are the $\binom{n}{2}$ inner products $\langle x_i, x_{i'} \rangle$ between all pairs of training observations.

It turns out that most of the $\hat{\alpha}_i$ can be zero:

$$f(x) = \beta_0 + \sum_{i \in \mathcal{S}} \hat{\alpha}_i \langle x, x_i \rangle$$

\mathcal{S} is the *support set* of indices i such that $\hat{\alpha}_i > 0$. [see slide 8]

Kernels and Support Vector Machines

- If we can compute inner-products between observations, we can fit a SV classifier. Can be quite abstract!

Kernels and Support Vector Machines

- If we can compute inner-products between observations, we can fit a SV classifier. Can be quite abstract!
- Some special *kernel functions* can do this for us. E.g.

$$K(x_i, x_{i'}) = \left(1 + \sum_{j=1}^p x_{ij} x_{i'j} \right)^d$$

computes the inner-products needed for d dimensional polynomials — $\binom{p+d}{d}$ basis functions!

Kernels and Support Vector Machines

- If we can compute inner-products between observations, we can fit a SV classifier. Can be quite abstract!
- Some special *kernel functions* can do this for us. E.g.

$$K(x_i, x_{i'}) = \left(1 + \sum_{j=1}^p x_{ij} x_{i'j} \right)^d$$

computes the inner-products needed for d dimensional polynomials — $\binom{p+d}{d}$ basis functions!

Try it for $p = 2$ and $d = 2$.

Kernel Functions and Support Vector Machines

can view a kernel conceptually as a way of measuring similarities (inner product) over expanded feature space

compute inner-products between observations, we can use a kernel classifier. Can be

- Some special *kernel functions* can

us solve the nonlinear classifier quicker?
- earlier SVC optimization with expanded features required $O(d^p)$ variables for optimization - many variables!
- in this kernel representation, we reduce the number of optimization variables to $n+1$ variables - much less!

Solving this particular kernel representation of the SVC (with a given kernel) - support vector machines (with kernel K)

$$K(x_i, x_{i'}) = \left(1 + \sum_{j=1}^p \sum_{k=1}^d x_{ij} x_{i'k} \right)^2$$

inner-products n-dimensional $(p+d)$ basis functions

how did we get this computational advantage? the kernel encodes all of the nonlinear features that we want, into a single similarity function.

Try it for $p = 2$ and $d = 2$.

- The solution has the form

$$f(x) = \beta_0 + \sum_{i \in S} \hat{\alpha}_i K(x, x_i).$$

view this as a way of
measuring similarities
between two features


$$K(x_i, x_{i'}) = \exp(-\gamma \sum_{j=1}^p (x_{ij} - x_{i'j})^2).$$

this kernel corresponds to
an infinite number of
nonlinear features
- Taylor expansion of K

$f(x) = +M$

$f(x) = 0$

$\hat{\alpha}_i K(x, x_i)$
 $i \in S$

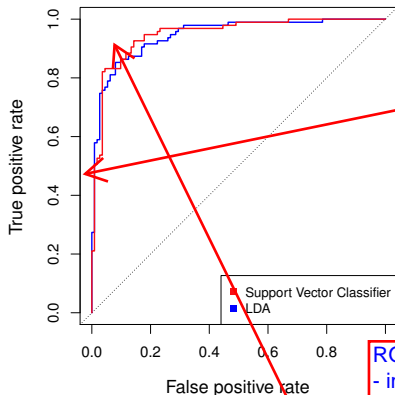


$f(x) = -M$

Implicit feature space;
very high dimensional.

Controls variance by
squashing down most
dimensions severely

Example: Heart Data



ROC curve is obtained by choosing t in $\hat{f}(X) > t$, and recording the true positive and false positive rates as t varies. Here we see

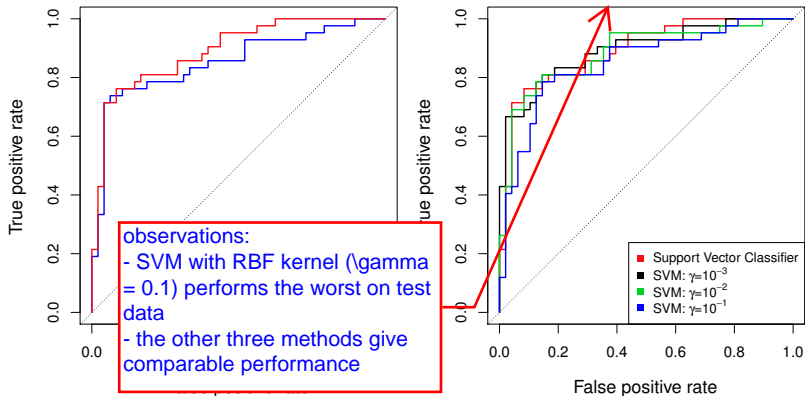
- sensitivity of classifier: for a person WITH heart disease, rate at which the classifier predicts heart disease (the larger the better)
- specificity of a classifier: for a person WITHOUT heart disease, rate at which the classifier predicts no heart disease (the larger the better)

on the training data: SVMs with radial kernel $\gamma = 0.1$ perform the best on training data

ROC (receiver operating curve):

- investigates diagnostic ability of binary classifier
- plot sensitivity vs. $1 - \text{specificity}$
- ROC curve is obtained by varying the discrimination threshold T in the classifier $f(x) = T$
 - $T = 0$: original classifier (from SVM)
 - T very large: all points are classified as 0 (perfect sensitivity)
 - T very negative: all points classified as 1 (perfect specificity)

Example continued: Heart Test Data



SVMs: more than 2 classes?

The SVM as defined works for $K = 2$ classes. What do we do if we have $K > 2$ classes?

SVMs: more than 2 classes?

The SVM as defined works for $K = 2$ classes. What do we do if we have $K > 2$ classes?

OVA One versus All. Fit K different 2-class SVM classifiers $\hat{f}_k(x)$, $k = 1, \dots, K$; each class versus the rest. Classify x^* to the class for which $\hat{f}_k(x^*)$ is largest.

SVMs: more than 2 classes?

The SVM as defined
we have $K > 2$ classes

- for each category k , we fit SVM classifiers with category k coded as +1, all others coded as -1
- recall that $|\hat{f}_k(x)|$ quantifies the distance from the point x to the classification hyperplane
- a larger value of this then suggests greater "confidence" that x belongs in the k -th class

to if

OVA One versus All. Fit K different 2-class SVM classifiers $\hat{f}_k(x)$, $k = 1, \dots, K$; each class versus the rest. Classify x^* to the class for which $\hat{f}_k(x^*)$ is largest.

OVO One versus One. Fit all $\binom{K}{2}$ pairwise classifiers $\hat{f}_{kl}(x)$. Classify x^* to the class that wins the most pairwise competitions.

- for each pair of categories (k,l) , fit SVM classifiers with class k coded as +1, class l coded as -1
- pick the most frequently assigned class

SVMs: more than 2 classes?

The SVM as defined works for $K = 2$ classes. What do we do if we have $K > 2$ classes?

OVA One versus All. Fit K different 2-class SVM classifiers $\hat{f}_k(x)$, $k = 1, \dots, K$; each class versus the rest. Classify x^* to the class for which $\hat{f}_k(x^*)$ is largest.


OVO One versus One. Fit all $\binom{K}{2}$ pairwise classifiers $\hat{f}_{k\ell}(x)$. Classify x^* to the class that wins the most pairwise competitions.

Which to choose? If K is not too large, use OVO.

Support Vector versus Logistic Regression?

With $f(X) = \beta_0 + \beta_1 X_1 + \dots + \beta_p X_p$ can rephrase support-vector classifier optimization as

$$\underset{\beta_0, \beta_1, \dots, \beta_p}{\text{minimize}} \left\{ \sum_{i=1}^n \max[0, 1 - y_i f(x_i)] + \lambda \sum_{j=1}^p \beta_j^2 \right\}$$



This has the form

loss plus penalty.

The loss is known as the *hinge loss*.

Very similar to “loss” in logistic regression (negative log-likelihood).

Which to use: SVM or Logistic Regression

In general, SVMs vs LR:

- SVMs marginally better for prediction accuracy
- LR: yields a probabilistic classifier which gives a measure of predictive (classification) uncertainty
 - should be considered when such uncertainty is important in the application

- When classes are (1) LR. So does LDA.
- When not, LR (with ridge penalty) and SVM very similar.
- If you wish to estimate probabilities, LR is the choice.
- For nonlinear boundaries, kernel SVMs are popular. Can use kernels with LR and LDA as well, but computations are more expensive.